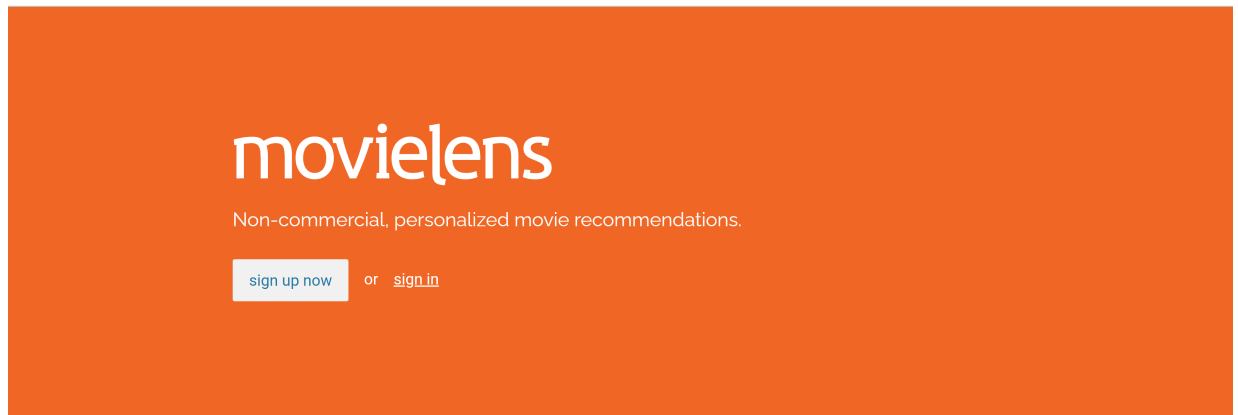


MovieLens Recommendation System

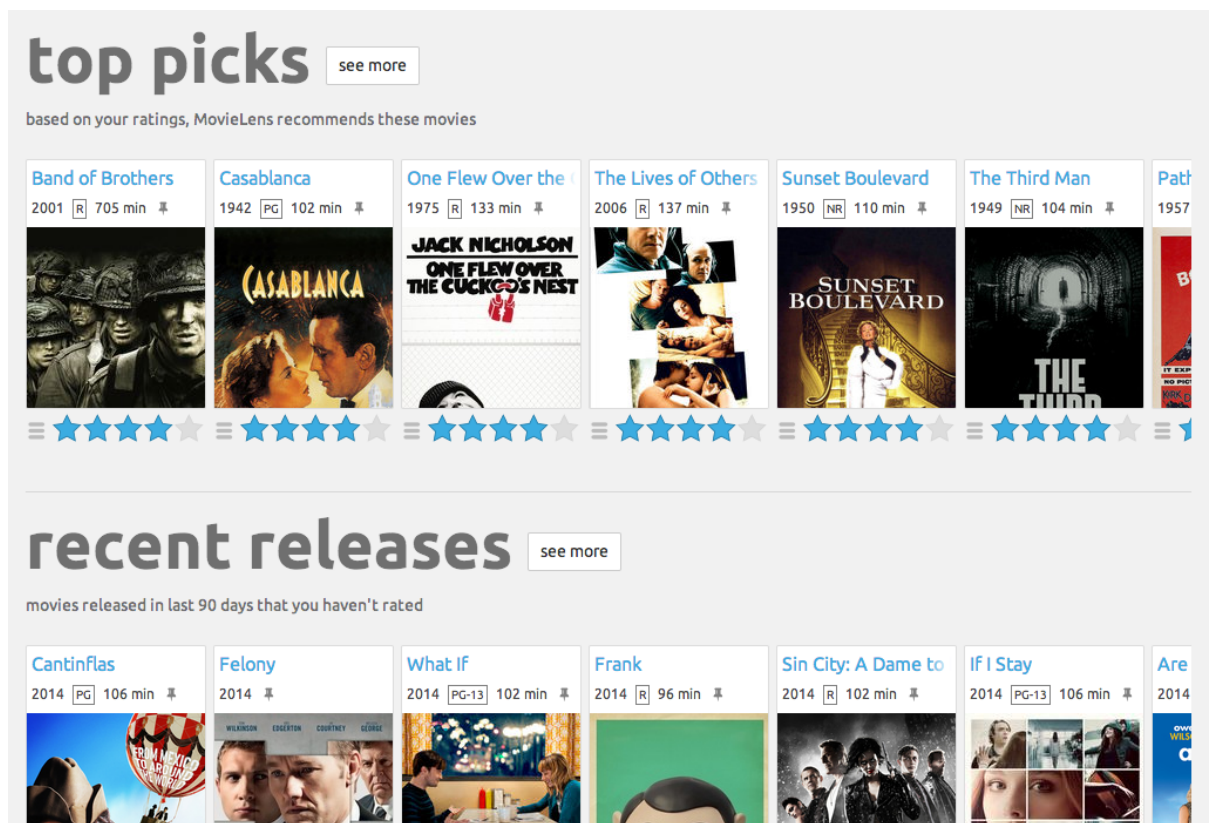
- Student name: Gamze Turan
- Student pace: self paced
- Scheduled project review date/time:
- Instructor name: Claude Fried
- Blog post URL: <https://ginaturan.blogspot.com/2022/08/natural-language-processing-nlp.html>
(<https://ginaturan.blogspot.com/2022/08/natural-language-processing-nlp.html>)



Overview

MovieLens is a website that helps users find movies they will like. It uses ratings given by the user to build a custom taste profile of that particular user and then utilizes that information to recommend other movies for the user to watch.

Business Understanding



Our goal is to build a variety of recommendation engines and improve upon predictions iteratively so that the end user can be provided with better movie suggestions.

Data Understanding

The datasets describe ratings and free-text tagging activities from [MovieLens](https://movielens.org/) (<https://movielens.org/>), a movie recommendation service.

Source: F. Maxwell Harper and Joseph A. Konstan. 2015. The MovieLens Datasets: History and Context. ACM Transactions on Interactive Intelligent Systems (TiiS) 5, 4: 19:1–19:19. <https://doi.org/10.1145/2827872> (<https://doi.org/10.1145/2827872>)

It contains 100836 ratings and 3683 tag applications across 9742 movies. These data were created by 610 users.

The dataset is distributed among four csv files: `links.csv` , `movies.csv` , `ratings.csv` , `tags.csv` .

```
In [1]: 1 # importing necessary libraries
2 import pandas as pd
3 import numpy as np
4 from scipy.sparse import csr_matrix
5 from sklearn.neighbors import NearestNeighbors
6
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9
10 import warnings
11 warnings.filterwarnings('ignore')
12
13 from surprise import Dataset
14 from surprise import Reader
15 from surprise import SVD, SVDpp
16 from surprise.prediction_algorithms import KNNWithMeans, KNNBasic, KNNWeighted
17 from surprise.model_selection import GridSearchCV
18 from surprise.model_selection import cross_validate
19
20 import ipywidgets as widgets
21 from ipywidgets import interact, interactive
22 from IPython.display import display, clear_output
```

Load data

Movies Data

Movie information is contained in the file movies.csv. Each line of this file after the header row represents one movie, and has the following format:

```
movieId,title,genres
```

- movieId : Unique id for each movie
- title : Name of movies followed by their year of release
- genres : categories that a movie might fall into separated by |

```
In [2]: 1 import pandas as pd
```

```
In [3]: 1 # movies data
2 movies_df = pd.read_csv('Data/movies.csv')
3 print('Size of movies data:', movies_df.shape)
4 movies_df.head()
```

Size of movies data: (9742, 3)

Out[3]:

| | movieId | title | genres |
|---|---------|------------------------------------|---|
| 0 | 1 | Toy Story (1995) | Adventure Animation Children Comedy Fantasy |
| 1 | 2 | Jumanji (1995) | Adventure Children Fantasy |
| 2 | 3 | Grumpier Old Men (1995) | Comedy Romance |
| 3 | 4 | Waiting to Exhale (1995) | Comedy Drama Romance |
| 4 | 5 | Father of the Bride Part II (1995) | Comedy |

```
In [4]: 1 movies_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9742 entries, 0 to 9741
Data columns (total 3 columns):
#   Column   Non-Null Count  Dtype
---  -
0   movieId  9742 non-null   int64
1   title     9742 non-null   object
2   genres    9742 non-null   object
dtypes: int64(1), object(2)
memory usage: 228.5+ KB
```

Some observations:

- There are no null values in the dataset and the datatypes of each of the columns are as they should be.
- movieID is consistent for all the other tables as well. So we can use this column to join together other datasets.
- We can extract the year of release for a movie, from the title.
- We will need to separate each genre into its own column to run meaningful analysis.

Ratings Data

Ratings information is contained in the file `ratings.csv`. Each line of this file after the header row represents one rating, and has the following format:

```
userId,movieId,rating,timestamp
```

- `userId` : Unique id for each user
- `movieId` : Unique id for each movie
- `rating` : Rating given by `userId` for `movieId` . Ratings are made on a 5-star scale with 0.5 increments.
- `timestamp` : Time when rating was given

```
In [5]: 1 #ratings data
        2 ratings_df = pd.read_csv('Data/ratings.csv')
        3 print('Size of ratings data:', ratings_df.shape)
        4 ratings_df.head()
```

Size of ratings data: (100836, 4)

Out[5]:

| | userId | movieId | rating | timestamp |
|---|--------|---------|--------|-----------|
| 0 | 1 | 1 | 4.0 | 964982703 |
| 1 | 1 | 3 | 4.0 | 964981247 |
| 2 | 1 | 6 | 4.0 | 964982224 |
| 3 | 1 | 47 | 5.0 | 964983815 |
| 4 | 1 | 50 | 5.0 | 964982931 |

```
In [6]: 1 ratings_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100836 entries, 0 to 100835
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   userId      100836 non-null  int64
1   movieId     100836 non-null  int64
2   rating      100836 non-null  float64
3   timestamp   100836 non-null  int64
dtypes: float64(1), int64(3)
memory usage: 3.1 MB
```

Some observations:

- There are no null values in the dataset.
- `timestamp` doesn't seem very useful for our current analysis and can be removed.

Links Data

The file `links.csv` contains identifiers that can be used to link this data to other data sources like IMDb. Each line of this file after the header row represents one imdb link, and has the following format:

```
movieId,imdbId,tmbId
```

- movieId : Unique id for each movie as used by <https://movielens.org> (<https://movielens.org>).
- imdbId : Unique id for each movie as used by <http://www.imdb.com> (<http://www.imdb.com>).
- tmbId : Unique id for each movie as used by <https://www.themoviedb.org> (<https://www.themoviedb.org>).

```
In [7]: 1 # Links data
2 links_df = pd.read_csv('Data/links.csv')
3 print('Size of links data:', links_df.shape)
4 links_df.head()
```

Size of links data: (9742, 3)

Out[7]:

| | movieId | imdbId | tmbId |
|---|---------|--------|---------|
| 0 | 1 | 114709 | 862.0 |
| 1 | 2 | 113497 | 8844.0 |
| 2 | 3 | 113228 | 15602.0 |
| 3 | 4 | 114885 | 31357.0 |
| 4 | 5 | 113041 | 11862.0 |

```
In [8]: 1 links_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9742 entries, 0 to 9741
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   movieId     9742 non-null   int64
1   imdbId      9742 non-null   int64
2   tmbId       9734 non-null   float64
dtypes: float64(1), int64(2)
memory usage: 228.5 KB
```

This table is not very useful for our current analyses and can be ignored.

Tags Data

Information regarding tags is contained in the file `tags.csv`. Each line of this file after the header row represents one tag applied to one movie by one user, and has the following format:

```
userId,movieId,tag,timestamp
```

- `userId` : Unique id for each user
- `movieId` : Unique id for each movie
- `tag` : User-generated metadata about the movie in forms of short meaningful phrases
- `timestamp` : Time when tag was provided by user

```
In [9]: 1 # tags data
        2 tags_df = pd.read_csv('Data/tags.csv')
        3 print('Size of tags data:', tags_df.shape)
        4 tags_df.head()
```

Size of tags data: (3683, 4)

Out[9]:

| | userId | movieId | tag | timestamp |
|---|--------|---------|-----------------|------------|
| 0 | 2 | 60756 | funny | 1445714994 |
| 1 | 2 | 60756 | Highly quotable | 1445714996 |
| 2 | 2 | 60756 | will ferrell | 1445714992 |
| 3 | 2 | 89774 | Boxing story | 1445715207 |
| 4 | 2 | 89774 | MMA | 1445715200 |

```
In [10]: 1 tags_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3683 entries, 0 to 3682
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   userId      3683 non-null   int64
1   movieId     3683 non-null   int64
2   tag         3683 non-null   object
3   timestamp   3683 non-null   int64
dtypes: int64(3), object(1)
memory usage: 115.2+ KB
```

EDA - Exploratory Data Analysis and Data Cleaning

```
In [11]: 1 # removing timestamp column from both ratings and tags as it
          2 ratings_df.drop(columns='timestamp', inplace=True)
          3 tags_df.drop(columns='timestamp', inplace=True)
```

```
In [12]: 1 ratings_df['rating'].describe()
```

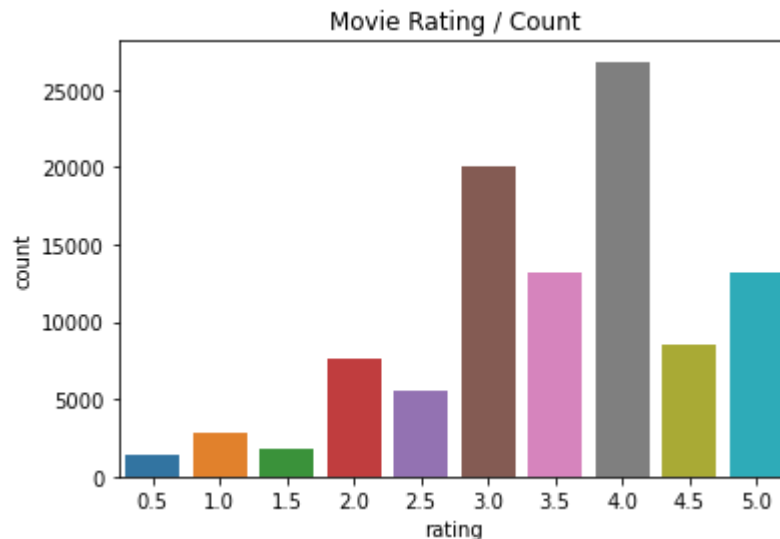
```
Out[12]: count    100836.000000
          mean       3.501557
          std       1.042529
          min       0.500000
          25%       3.000000
          50%       3.500000
          75%       4.000000
          max       5.000000
          Name: rating, dtype: float64
```

```
In [13]: 1 ratings_df['rating'].value_counts()
```

```
Out[13]: 4.0    26818
          3.0    20047
          5.0    13211
          3.5    13136
          4.5     8551
          2.0     7551
          2.5     5550
          1.0     2811
          1.5     1791
          0.5     1370
          Name: rating, dtype: int64
```



```
In [14]: 1 sns.countplot(ratings_df['rating'])
2 plt.title('Movie Rating / Count');
3 #plt.savefig('Images/movie_rating')
```



I observe that the mean rating given by users is approximately 3.5 - 4 is the most common rating in the dataset. Most of the ratings in the dataset are above 3.

Now, I will extract movie release year from movie title.

```
In [15]: 1 # Extracting release year from movie title
2 movies_df['year'] = movies_df['title'].str.extract('.*\((.*)\)ate.*', expand=True)
```

```
In [16]: 1 movies_df['year'].unique()
```

```
Out[16]: array(['1995', '1994', '1996', '1976', '1992', '1967', '1993', '1964',
'1977', '1965', '1982', '1990', '1991', '1989', '1937', '1940',
'1969', '1981', '1973', '1970', '1955', '1959', '1968', '1988',
'1997', '1972', '1943', '1952', '1951', '1957', '1961', '1958',
'1954', '1934', '1944', '1960', '1963', '1942', '1941', '1953',
'1939', '1950', '1946', '1945', '1938', '1947', '1935', '1936',
'1956', '1949', '1932', '1975', '1974', '1971', '1979', '1987',
'1986', '1980', '1978', '1985', '1966', '1962', '1983', '1984',
'1948', '1933', '1931', '1922', '1998', '1929', '1930', '1927',
'1928', '1999', '2000', '1926', '1919', '1921', '1925', '1923',
'2001', '2002', '2003', '1920', '1915', '1924', '2004', '1916',
'1917', '2005', '2006', '1902', nan, '1903', '2007', '2008',
'2009', '2010', '2011', '2012', '2013', '2014', '2015', '2016',
'2017', '2018', '1908', '2006-2007'], dtype=object)
```

One value that is slightly off is '2006-2007' and the other is NaN. Let's check which movies do they correspond to:

```
In [17]: 1 movies_df[movies_df['year'] == "2006-2007"]
```

Out[17]:

| | movieid | title | genres | year |
|------|---------|-----------------------------------|--------------------|-----------|
| 9518 | 171749 | Death Note: Desu nôto (2006–2007) | (no genres listed) | 2006–2007 |

```
In [18]: 1 # Changing this to 2007
2 movies_df['year'] = movies_df['year'].replace("2006-2007", "2007")
```

```
In [19]: 1 # movies with no year information
2 movies_df[pd.isna(movies_df['year'])]
```

Out[19]:

| | movieid | title | genres | year |
|------|---------|---|------------------------|------|
| 6059 | 40697 | Babylon 5 | Sci-Fi | NaN |
| 9031 | 140956 | Ready Player One | Action Sci-Fi Thriller | NaN |
| 9091 | 143410 | Hyena Road | (no genres listed) | NaN |
| 9138 | 147250 | The Adventures of Sherlock Holmes and Doctor W... | (no genres listed) | NaN |
| 9179 | 149334 | Nocturnal Animals | Drama Thriller | NaN |
| 9259 | 156605 | Paterson | (no genres listed) | NaN |
| 9367 | 162414 | Moonlight | Drama | NaN |
| 9448 | 167570 | The OA | (no genres listed) | NaN |
| 9514 | 171495 | Cosmos | (no genres listed) | NaN |
| 9515 | 171631 | Maria Bamford: Old Baby | (no genres listed) | NaN |
| 9525 | 171891 | Generation Iron 2 | (no genres listed) | NaN |
| 9611 | 176601 | Black Mirror | (no genres listed) | NaN |

```
In [20]: 1 # As many of these don't have any genres as well, we will drop these rows
2 movies_df = movies_df.dropna(subset=['year'], how='any')
```

```
In [21]: 1 movies_df['year'] = movies_df['year'].astype(int)
```

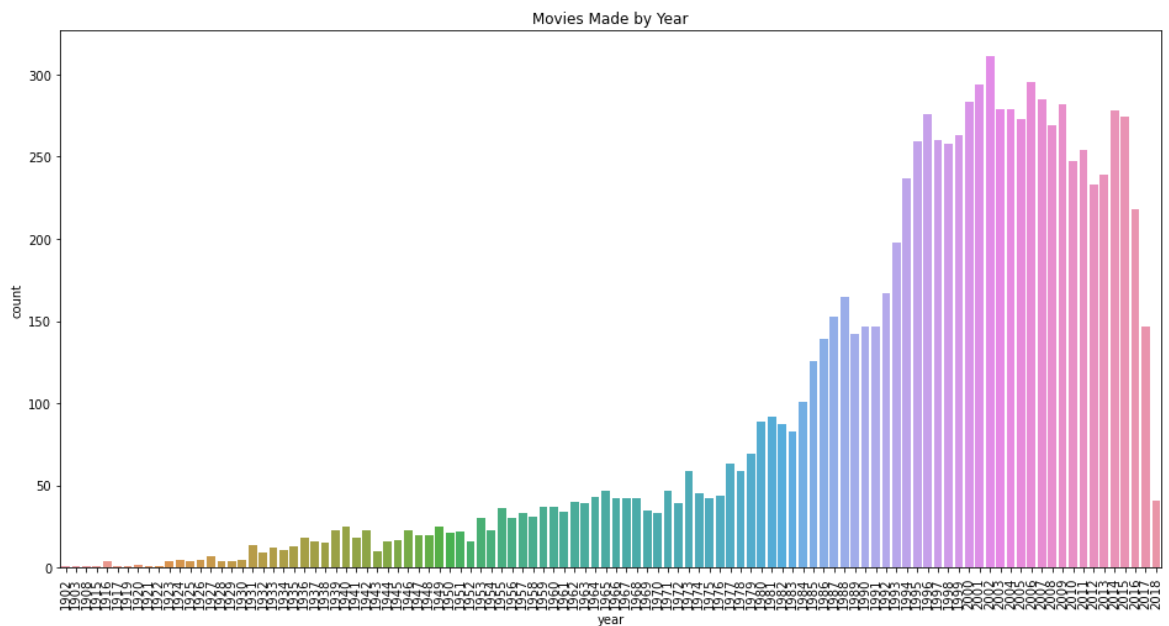
In [22]: 1 movies_df['year'].describe()

```
Out[22]: count    9730.000000
         mean    1994.614902
         std      18.534692
         min    1902.000000
         25%    1988.000000
         50%    1999.000000
         75%    2008.000000
         max    2018.000000
         Name: year, dtype: float64
```

In [23]: 1 movies_df['year'].value_counts()

```
Out[23]: 2002    311
         2006    295
         2001    294
         2007    285
         2000    283
         ...
         1917     1
         1902     1
         1903     1
         1919     1
         1908     1
         Name: year, Length: 106, dtype: int64
```

```
In [24]: 1 # number of movie in year
2 plt.figure(figsize=(16,8))
3 sns.countplot(movies_df['year'])
4 plt.xticks(rotation=90)
5 plt.title('Movies Made by Year');
6 #plt.savefig('Images/movie_year')
```



In this dataset, we have movies starting as early as 1902 and the latest movie is from 2018. The year with the maximum number of movies in this dataset is 2002 with 311 movies.

Now, let's look for any duplicate values.

```
In [25]: 1 # Validation check
2 print('Is there a duplicate value in a column movieId? Ans:',not movies_
3 print('Is there a duplicate value in a column title? Ans:',not movies_df
4
5 movies_df[movies_df.duplicated(["title"], keep=False)]
```

Is there a duplicate value in a column movieId? Ans: False

Is there a duplicate value in a column title? Ans: True

Out[25]:

| | movieId | | title | genres | year |
|------|---------|---|-----------------|-------------------------------------|------|
| 650 | 838 | | Emma (1996) | Comedy Drama Romance | 1996 |
| 2141 | 2851 | | Saturn 3 (1980) | Adventure Sci-Fi Thriller | 1980 |
| 4169 | 6003 | Confessions of a Dangerous Mind (2002) | | Comedy Crime Drama Thriller | 2002 |
| 5601 | 26958 | | Emma (1996) | Romance | 1996 |
| 5854 | 32600 | | Eros (2004) | Drama | 2004 |
| 5931 | 34048 | War of the Worlds (2005) | | Action Adventure Sci-Fi Thriller | 2005 |
| 6932 | 64997 | War of the Worlds (2005) | | Action Sci-Fi | 2005 |
| 9106 | 144606 | Confessions of a Dangerous Mind (2002) | | Comedy Crime Drama Romance Thriller | 2002 |
| 9135 | 147002 | | Eros (2004) | Drama Romance | 2004 |
| 9468 | 168358 | | Saturn 3 (1980) | Sci-Fi Thriller | 1980 |

There are 5 movies which are duplicated in our dataset. They have the same title but different movieId. The problem is that the same movieId is then used in other tables as well. Let's replace the movieIds in each table.

```
In [26]: 1 movie_id_change = {838:26958, 2851:168358, 6003:144606,32600:147002,34048:
2 movies_df['movieId'].replace(movie_id_change,inplace=True)
3 movies_df = movies_df.drop_duplicates(subset=["movieId","title"])
4
5 ratings_df['movieId'].replace(movie_id_change,inplace=True)
6 tags_df['movieId'].replace(movie_id_change,inplace=True)
7 links_df['movieId'].replace(movie_id_change,inplace=True)
```

```
In [27]: 1 # Validation check
2 print('Is there a duplicate value in a column movieId? Ans:',not movies_
3 print('Is there a duplicate value in a column title? Ans:',not movies_df
```

Is there a duplicate value in a column movieId? Ans: False

Is there a duplicate value in a column title? Ans: False

Now let's extract individual genres from the genres column. But, first we need to replace a value in

this column. Null value in genres column is given as (no genres listed) .

```
In [28]: 1 import numpy as np
```

```
In [29]: 1 movies_df['genres'] = movies_df['genres'].replace('(no genres listed)',  
2 print('Number of missing values in genres column:',movies_df['genres'].i  
3  
4 # dropping rows with missing genres  
5 movies_df = movies_df.dropna(subset=['genres'],how='any')  
6 movies_df = movies_df.reset_index(drop=True)
```

Number of missing values in genres column: 26

Now, I will use One-Hot Encoding and create columns for each genre

```

In [30]: 1 # Serepate the Genres Column and Encoding them with One-Hot Encoding
2 genres = []
3 for i in range(len(movies_df.genres)):
4     for x in movies_df.genres[i].split('|'):
5         if x not in genres:
6             genres.append(x)
7
8 len(genres)
9 for x in genres:
10     movies_df[x] = 0
11 for i in range(len(movies_df.genres)):
12     for x in movies_df.genres[i].split('|'):
13         movies_df[x][i]=1
14
15 #dropping the genres column as it's a no longer required
16 movies_df.drop(columns='genres', inplace=True)
17 movies_df.sort_index(inplace=True)
18 movies_df

```

Out[30]:

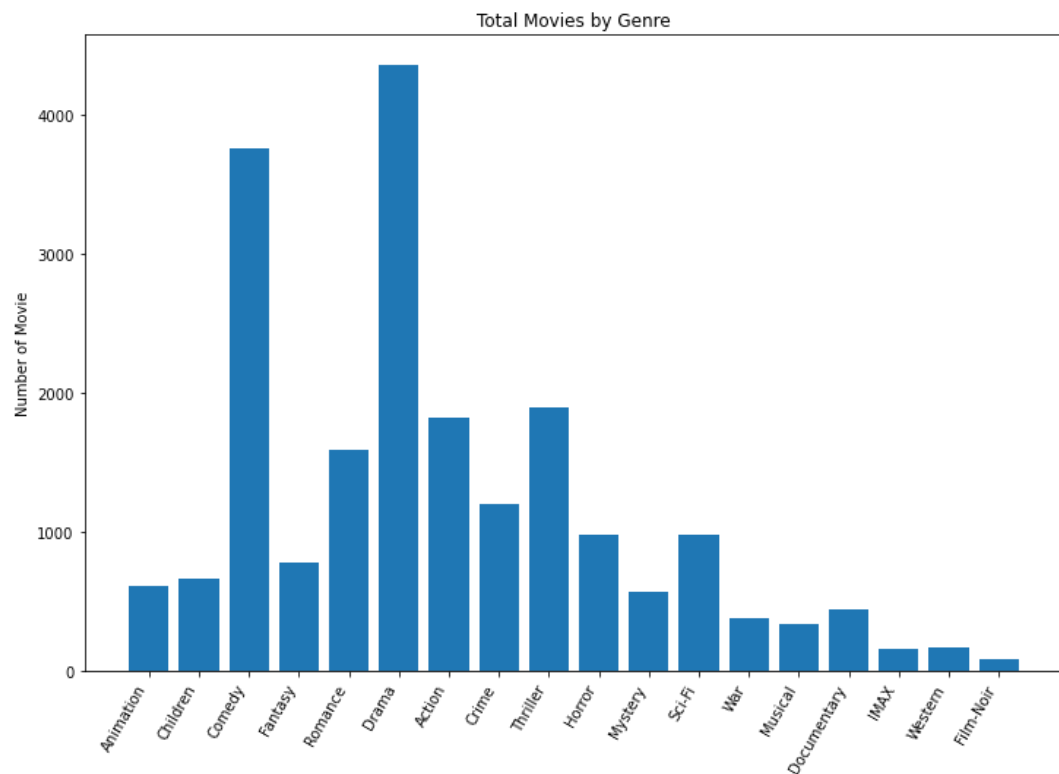
| | movieid | title | year | Adventure | Animation | Children | Comedy | Fantasy | Ro |
|------|---------|---|------|-----------|-----------|----------|--------|---------|----|
| 0 | 1 | Toy Story (1995) | 1995 | 1 | 1 | 1 | 1 | 1 | |
| 1 | 2 | Jumanji (1995) | 1995 | 1 | 0 | 1 | 0 | 1 | |
| 2 | 3 | Grumpier Old Men (1995) | 1995 | 0 | 0 | 0 | 1 | 0 | |
| 3 | 4 | Waiting to Exhale (1995) | 1995 | 0 | 0 | 0 | 1 | 0 | |
| 4 | 5 | Father of the Bride Part II (1995) | 1995 | 0 | 0 | 0 | 1 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 9694 | 193581 | Black Butler: Book of the Atlantic (2017) | 2017 | 0 | 1 | 0 | 1 | 1 | |
| 9695 | 193583 | No Game No Life: Zero (2017) | 2017 | 0 | 1 | 0 | 1 | 1 | |
| 9696 | 193585 | Flint (2017) | 2017 | 0 | 0 | 0 | 0 | 0 | |

| | movieId | title | year | Adventure | Animation | Children | Comedy | Fantasy | R |
|------|---------|-------------------------------------|------|-----------|-----------|----------|--------|---------|---|
| 9697 | 193587 | Bungo Stray Dogs: Dead Apple (2018) | 2018 | 0 | 1 | 0 | 0 | 0 | |
| 9698 | 193609 | Andrew Dice Clay: Dice Rules (1991) | 1991 | 0 | 0 | 0 | 1 | 0 | |

9699 rows × 22 columns


```
In [31]: 1 # plotting genres popularity
2 x = {}
3 for i in movies_df.columns[4:23]:
4     x[i] = movies_df[i].sum()
5     print(f"{i:<15}{x[i]:>10}")
6
7 plt.figure(figsize=(12,8))
8 plt.bar(height = x.values(), x=x.keys())
9 plt.xticks(rotation=60, ha='right')
10 plt.ylabel('Number of Movie')
11 plt.title('Total Movies by Genre')
12 #plt.savefig('Images/genres')
13 plt.show()
```

| | |
|-------------|------|
| Animation | 611 |
| Children | 664 |
| Comedy | 3755 |
| Fantasy | 779 |
| Romance | 1593 |
| Drama | 4357 |
| Action | 1826 |
| Crime | 1198 |
| Thriller | 1890 |
| Horror | 978 |
| Mystery | 573 |
| Sci-Fi | 976 |
| War | 382 |
| Musical | 334 |
| Documentary | 440 |
| IMAX | 158 |
| Western | 167 |
| Film-Noir | 87 |



Drama is the most popular genre with 4357 movies, followed by Comedy with 3755 movies.

I will now merge the ratings and movie dataframe to get the average rating and num of ratings per movie.

```
In [32]: 1 # movies with no ratings receive 0 as their average rating
2 mean_rating = ratings_df.groupby('movieId').rating.mean().rename('mean r
3 num_rating = ratings_df.groupby('movieId').userId.count().rename('num ra
4
5 movies_df = pd.merge(movies_df, mean_rating, how='left', on='movieId')
6 movies_df = pd.merge(movies_df, num_rating, how='left', on='movieId')
7
8 movies_df['mean rating'].fillna(0, inplace=True)
9 movies_df['num rating'].fillna(0, inplace=True)
10
11 movies_df[['title', 'mean rating', 'num rating']]
```

Out[32]:

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 0 | Toy Story (1995) | 3.920930 | 215.0 |
| 1 | Jumanji (1995) | 3.431818 | 110.0 |
| 2 | Grumpier Old Men (1995) | 3.259615 | 52.0 |
| 3 | Waiting to Exhale (1995) | 2.357143 | 7.0 |
| 4 | Father of the Bride Part II (1995) | 3.071429 | 49.0 |
| ... | ... | ... | ... |
| 9694 | Black Butler: Book of the Atlantic (2017) | 4.000000 | 1.0 |
| 9695 | No Game No Life: Zero (2017) | 3.500000 | 1.0 |
| 9696 | Flint (2017) | 3.500000 | 1.0 |
| 9697 | Bungo Stray Dogs: Dead Apple (2018) | 3.500000 | 1.0 |
| 9698 | Andrew Dice Clay: Dice Rules (1991) | 4.000000 | 1.0 |

9699 rows × 3 columns

Adding mean rating and num rating columns to our dataset would allow us to understand which movie is well loved or reviewed in our database. I will make use of this information in the following sections.

Naive Recommendation Engine

This system, would use overall ratings and genres to recommend movies. This could work well in helping resolve the cold-start problem as well in the future.

As the first initial model, I can recommend the top 10 most popular movies(movies with most number of ratings) in our database to a new user.

```
In [33]: 1 movie_ratings = movies_df[['title', 'mean rating', 'num rating', ]]
```

```
In [34]: 1 movie_ratings.sort_values(by=['num rating'], ascending=False).head(10)
```

Out[34]:

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 314 | Forrest Gump (1994) | 4.164134 | 329.0 |
| 277 | Shawshank Redemption, The (1994) | 4.429022 | 317.0 |
| 257 | Pulp Fiction (1994) | 4.197068 | 307.0 |
| 510 | Silence of the Lambs, The (1991) | 4.161290 | 279.0 |
| 1939 | Matrix, The (1999) | 4.192446 | 278.0 |
| 224 | Star Wars: Episode IV - A New Hope (1977) | 4.231076 | 251.0 |
| 418 | Jurassic Park (1993) | 3.750000 | 238.0 |
| 97 | Braveheart (1995) | 4.031646 | 237.0 |
| 507 | Terminator 2: Judgment Day (1991) | 3.970982 | 224.0 |
| 461 | Schindler's List (1993) | 4.225000 | 220.0 |

As expected, all listed movies are internationally acclaimed hollywood classics.

Now let's look at the top movies with the highest ratings in our database

```
In [35]: 1 # top 10 movies by its mean ratings
2 movie_ratings.sort_values(by=['mean rating'], ascending=False).head(10)
```

Out[35]:

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 7598 | Idiots and Angels (2008) | 5.0 | 1.0 |
| 8670 | Stuart Little 3: Call of the Wild (2005) | 5.0 | 1.0 |
| 3110 | Reform School Girls (1986) | 5.0 | 1.0 |
| 8501 | One I Love, The (2014) | 5.0 | 1.0 |
| 8513 | Laggies (2014) | 5.0 | 1.0 |
| 3081 | Monster Squad, The (1987) | 5.0 | 1.0 |
| 8547 | Crippled Avengers (Can que) (Return of the 5 D... | 5.0 | 1.0 |
| 3067 | Hollywood Shuffle (1987) | 5.0 | 1.0 |
| 8587 | Watermark (2014) | 5.0 | 1.0 |
| 8606 | Hellbenders (2012) | 5.0 | 1.0 |

While these movies are rated quite high, they are not popular and only have 1 rating. This is not reliable information. I need to set a threshold of the minimum number of ratings a movie must have

```
In [36]: 1 #creating minimum number of ratings
2 minimum_num_ratings = 200
3 movie_ratings[movie_ratings['num rating']>minimum_num_ratings].sort_valu
```

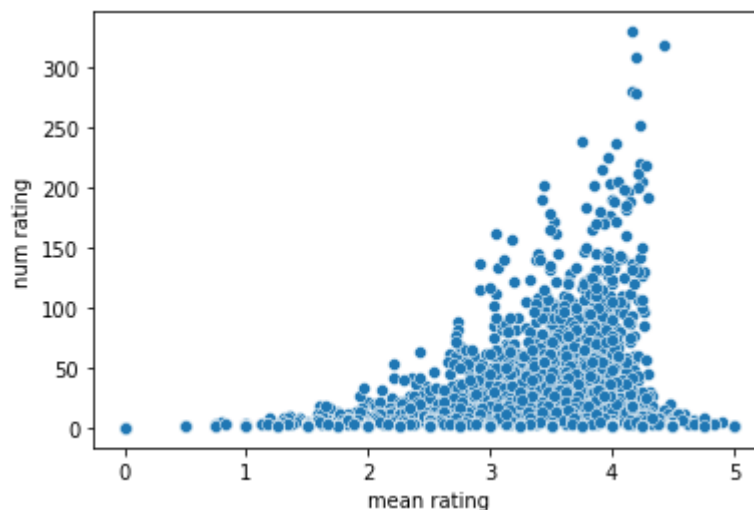
Out[36]:

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 277 | Shawshank Redemption, The (1994) | 4.429022 | 317.0 |
| 2226 | Fight Club (1999) | 4.272936 | 218.0 |
| 46 | Usual Suspects, The (1995) | 4.237745 | 204.0 |
| 224 | Star Wars: Episode IV - A New Hope (1977) | 4.231076 | 251.0 |
| 461 | Schindler's List (1993) | 4.225000 | 220.0 |
| 898 | Star Wars: Episode V - The Empire Strikes Back... | 4.215640 | 211.0 |
| 257 | Pulp Fiction (1994) | 4.197068 | 307.0 |
| 1939 | Matrix, The (1999) | 4.192446 | 278.0 |
| 314 | Forrest Gump (1994) | 4.164134 | 329.0 |
| 510 | Silence of the Lambs, The (1991) | 4.161290 | 279.0 |

Now I see movies which are popular as well as loved by the users.

Also, let's see if there is correlation between movies with higher number of ratings and movies with high average rating.

```
In [37]: 1 # mean rating and total number of rating scatterplot
2 sns.scatterplot(data=movies_df, x='mean rating', y='num rating');
3 #plt.savefig('Images/mean_num_rating')
```



As expected, movies that are good also receive more ratings.

I can also use Genre information to further refine our recommendations.

Suppose, there is a user who wants recommendation for an action movie, then:

```
In [38]: 1 # checking users that interest on action movie
          2 user_genre = 'Action'
          3 movie_ratings[movies_df[user_genre] == 1].sort_values(by=['mean rating'])
```

Out[38]:

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 8547 | Crippled Avengers (Can que) (Return of the 5 D... | 5.0 | 1.0 |
| 7488 | Faster (2010) | 5.0 | 1.0 |
| 8145 | Justice League: Doom (2012) | 5.0 | 1.0 |
| 8973 | Tokyo Tribe (2014) | 5.0 | 1.0 |
| 3759 | Shogun Assassin (1980) | 5.0 | 1.0 |
| 3908 | The Big Bus (1976) | 5.0 | 1.0 |
| 3110 | Reform School Girls (1986) | 5.0 | 1.0 |
| 7900 | Superman/Batman: Public Enemies (2009) | 5.0 | 1.0 |
| 1647 | Knock Off (1998) | 5.0 | 1.0 |
| 4045 | Galaxy of Terror (Quest) (1981) | 5.0 | 1.0 |

Again setting a threshold on the minimum number of ratings

```
In [39]: 1 # setting minimum number of rating on users that interest on action movie
          2 user_genre = 'Action'
          3 minimum_num_ratings = 200
          4 movie_ratings[(movies_df[user_genre] == 1) & (movies_df['num rating'] > mi
```

Out[39]:

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 2226 | Fight Club (1999) | 4.272936 | 218.0 |
| 224 | Star Wars: Episode IV - A New Hope (1977) | 4.231076 | 251.0 |
| 898 | Star Wars: Episode V - The Empire Strikes Back... | 4.215640 | 211.0 |
| 1939 | Matrix, The (1999) | 4.192446 | 278.0 |
| 97 | Braveheart (1995) | 4.031646 | 237.0 |
| 507 | Terminator 2: Judgment Day (1991) | 3.970982 | 224.0 |
| 418 | Jurassic Park (1993) | 3.750000 | 238.0 |
| 615 | Independence Day (a.k.a. ID4) (1996) | 3.445545 | 202.0 |

Naive Recommendation Engine (New User)

Let's combine all the techniques that we used above to build a basic recommendation engine:

1. If no information is available for the user, then recommend movies rated more than default threshold of 100 times with the highest ratings.
2. If user sets a threshold, then recommend movies rated more than threshold times with the highest ratings.
3. If genre is presented, then recommend movies from that genre rated more than threshold times with the highest ratings.

```

In [40]: ▶ 1 def naive_recommendation(threshold,fav_genre):
2
3     minimum_num_ratings = threshold
4     if fav_genre == 'All':
5         result = movie_ratings[(movies_df['num rating']>minimum_num_ratings)
6     else:
7         result = movie_ratings[(movies_df[fav_genre] == 1) & (movies_df['num rating']>minimum_num_ratings)]
8
9     print('\n\nThese are the recommendations for the users with the following filters:')
10    print('Minimum number of ratings:',threshold)
11    print("User's choice of genre:",fav_genre)
12    display(result)
13
14
15    genres = ['All',
16              'Animation',
17              'Children',
18              'Comedy',
19              'Fantasy',
20              'Romance',
21              'Drama',
22              'Action',
23              'Crime',
24              'Thriller',
25              'Horror',
26              'Mystery',
27              'Sci-Fi',
28              'War',
29              'Musical',
30              'Documentary',
31              'IMAX',
32              'Western',
33              'Film-Noir']
34
35    w = interactive(naive_recommendation, threshold=widgets.IntSlider(min=0, max=100, value=100),
36                   fav_genre=widgets.Dropdown(options=genres, description='Genre'))
37
38    display(w)

```

threshold  100

Genre

These are the recommendations for the users with the following filters
 Minimum number of ratings: 100
 User's choice of genre: All

| | title | mean rating | num rating |
|-----|----------------------------------|-------------|------------|
| 277 | Shawshank Redemption, The (1994) | 4.429022 | 317.0 |

| | title | mean rating | num rating |
|------|---|-------------|------------|
| 659 | Godfather, The (1972) | 4.289062 | 192.0 |
| 2226 | Fight Club (1999) | 4.272936 | 218.0 |
| 922 | Godfather: Part II, The (1974) | 4.259690 | 129.0 |
| 6313 | Departed, The (2006) | 4.252336 | 107.0 |
| 914 | Goodfellas (1990) | 4.250000 | 126.0 |
| 6708 | Dark Knight, The (2008) | 4.238255 | 149.0 |
| 46 | Usual Suspects, The (1995) | 4.237745 | 204.0 |
| 899 | Princess Bride, The (1987) | 4.232394 | 142.0 |
| 224 | Star Wars: Episode IV - A New Hope (1977) | 4.231076 | 251.0 |

Colloborative Filtering

Item based Using Correlation

I will now design a recomendation engine that uses the correlation between the ratings assined to different movies, in order to find the similarity between the movies.

Let's create a matrix where each column is a movie name and esch row contains the rating assigned by a specific user to that movie.

```
In [41]: 1 # merging ratings and movies data
2 df = pd.merge(ratings_df, movies_df, how='left', on = 'movieId')
3 # creating matrix
4 movie_user_matrix = df.pivot_table(index='userId', columns='title', values='rating')
5 movie_user_matrix
```

Out[41]:

| | title | '71 (2014) | 'Hellboy': The Seeds of Creation (2004) | 'Round Midnight (1986) | 'Salem's Lot (2004) | 'Til There Was You (1997) | 'Tis the Season for Love (2015) | 'burbs, The (1989) | 'night Mother (1986) | (500) Days of Summer (2009) | *batte inclu (19 |
|--------|-------|---------------|---|------------------------------|---------------------------|---------------------------------------|---|--------------------------|----------------------------|--------------------------------------|------------------------|
| userId | | | | | | | | | | | |
| 1 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 2 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 3 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 4 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 5 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 606 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 607 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 608 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 609 | | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 610 | | 4.0 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | 3.5 | NaN |

610 rows × 9681 columns



As most users will only review a few movies, majority of the values in this matrix will be NaN.

Now, let's consider that you wish look at movies that are similar to Lion King.

```
In [42]: 1 movie_name = "Lion King, The (1994)"
          2 movie_ratings_df = movie_user_matrix[movie_name]
          3 movie_ratings_df.head()
```

```
Out[42]: userId
1      NaN
2      NaN
3      NaN
4      NaN
5      3.0
Name: Lion King, The (1994), dtype: float64
```

Retrieving movies where the ratings are extremely correlated with Lion King:

```
In [43]: 1 correlation = movie_user_matrix.corrwith(movie_ratings_df)
          2 similar_movies = pd.DataFrame(correlation, columns=['Correlation'])
          3 # removing nulls
          4 similar_movies.dropna(inplace=True)
          5
          6 # Top 10 highly correlated movies with Lion King
          7 similar_movies.sort_values('Correlation', ascending=False).head(10)
```

```
Out[43]:
```

| | Correlation |
|-----------------------------------|-------------|
| Despicable Me 3 (2017) | 1.0 |
| Sound of Thunder, A (2005) | 1.0 |
| Damsels in Distress (2011) | 1.0 |
| Danny Deckchair (2003) | 1.0 |
| Joy (2015) | 1.0 |
| Only the Lonely (1991) | 1.0 |
| Soul Plane (2004) | 1.0 |
| Jobs (2013) | 1.0 |
| Ordet (Word, The) (1955) | 1.0 |
| Babe, The (1992) | 1.0 |

Some of these ratings make sense, but some don't. Let's bring number of ratings in picture here. I will set a threshold of 50 to the number of ratings here:

```
In [44]: 1 similar_movies = pd.merge(similar_movies, movies_df[['title','num rating
2 similar_movies.set_index('title', inplace=True)
3
4 # Top 10 highly correlated movies with Lion King that have been rated mo
5 threshold = 50
6 similar_movies.sort_values('Correlation', ascending=False)[similar_movie
```

Out[44]:

| | Correlation | num rating |
|--------------------------------|-------------|------------|
| title | | |
| Lion King, The (1994) | 1.000000 | 172.0 |
| Guardians of the Galaxy (2014) | 0.673887 | 59.0 |
| X2: X-Men United (2003) | 0.596938 | 76.0 |
| Aladdin (1992) | 0.591660 | 183.0 |
| While You Were Sleeping (1995) | 0.565303 | 98.0 |
| Casper (1995) | 0.555249 | 62.0 |
| Hook (1991) | 0.541501 | 53.0 |
| Grumpier Old Men (1995) | 0.541416 | 52.0 |
| Predator (1987) | 0.533184 | 61.0 |
| Beautiful Mind, A (2001) | 0.533109 | 123.0 |

While there are some movies like Aladdin that do make sense here, most don't. This means that maybe we need to improve our model or add more information to help the model.

User Based Collaborative Filtering using surprise library

For our next few models, we will utilize the `surprise` library which allows us to build complex recommendation engine pipelines effortlessly. I will try out a total of five algorithms.

First three are KNN based algorithms:

- KNNBasic
- KNNWithMeans
- KNNWithZone

Next two are matrix factorization based algorithms:

- SVD

- SVDpp

I will evaluate top models from each list and then do grid search on them to search for the best hyperparameters.

But first, I will convert our dataset into something that the surprise library can understand.

```
In [45]: 1 from surprise import Dataset
          2 from surprise import Reader
          3 from surprise import SVD, SVDpp
          4 from surprise.prediction_algorithms import KNNWithMeans, KNNBasic, KNNWith
          5 from surprise.model_selection import GridSearchCV
          6 from surprise.model_selection import cross_validate
```

```
In [46]: 1 # read in values as Surprise dataset
          2 reader = Reader(rating_scale=(0,5))
          3 data = Dataset.load_from_df(ratings_df, reader)
          4
          5 #generating a trainset
          6 dataset = data.build_full_trainset()
          7 print('Number of users: ', dataset.n_users, '\n')
          8 print('Number of items: ', dataset.n_items)
```

Number of users: 610

Number of items: 9719

```
In [47]: 1 # knn algorithms
2 cv_knn_basic = cross_validate(KNNBasic(), data, cv=5, n_jobs=5, verbose=
3 cv_knn_means = cross_validate(KNNWithMeans(), data, cv=5, n_jobs=5, verbo
4 cv_knn_z = cross_validate(KNNWithZScore(), data, cv=5, n_jobs=5, verbose=
```

Evaluating RMSE, MAE of algorithm KNNBasic on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|--------|
| RMSE (testset) | 0.9449 | 0.9448 | 0.9512 | 0.9452 | 0.9490 | 0.9470 | 0.0026 |
| MAE (testset) | 0.7224 | 0.7265 | 0.7287 | 0.7227 | 0.7268 | 0.7254 | 0.0025 |
| Fit time | 0.91 | 0.82 | 1.06 | 0.63 | 0.58 | 0.80 | 0.18 |
| Test time | 5.95 | 7.89 | 6.06 | 6.71 | 7.43 | 6.81 | 0.76 |

Evaluating RMSE, MAE of algorithm KNNWithMeans on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|--------|
| RMSE (testset) | 0.9030 | 0.8977 | 0.8899 | 0.8941 | 0.9009 | 0.8971 | 0.0047 |
| MAE (testset) | 0.6898 | 0.6844 | 0.6807 | 0.6852 | 0.6858 | 0.6852 | 0.0029 |
| Fit time | 0.65 | 0.71 | 0.79 | 0.84 | 0.88 | 0.77 | 0.08 |
| Test time | 7.56 | 8.83 | 9.13 | 9.08 | 8.43 | 8.61 | 0.58 |

Evaluating RMSE, MAE of algorithm KNNWithZScore on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|--------|
| RMSE (testset) | 0.8896 | 0.8925 | 0.9070 | 0.8933 | 0.8988 | 0.8963 | 0.0062 |
| MAE (testset) | 0.6792 | 0.6765 | 0.6865 | 0.6747 | 0.6821 | 0.6798 | 0.0042 |
| Fit time | 0.94 | 0.86 | 1.15 | 1.21 | 1.20 | 1.07 | 0.14 |
| Test time | 8.74 | 7.71 | 7.51 | 9.64 | 9.53 | 8.63 | 0.89 |

```
In [48]: 1 # matrix factorization algorithms
2 # run time SVDpp approx.20min
3 cv_svd = cross_validate(SVD(), data, cv=5, n_jobs=5, verbose=True)
4 cv_svd_pp = cross_validate(SVDpp(), data, cv=5, n_jobs=5, verbose=True)
```

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|--------|
| RMSE (testset) | 0.8740 | 0.8730 | 0.8694 | 0.8789 | 0.8776 | 0.8746 | 0.0034 |
| MAE (testset) | 0.6710 | 0.6732 | 0.6687 | 0.6736 | 0.6751 | 0.6723 | 0.0022 |
| Fit time | 19.34 | 24.84 | 19.97 | 22.00 | 22.66 | 21.76 | 1.97 |
| Test time | 0.50 | 0.43 | 0.52 | 0.59 | 0.57 | 0.52 | 0.06 |

Evaluating RMSE, MAE of algorithm SVDpp on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|---------|---------|---------|---------|---------|---------|--------|
| RMSE (testset) | 0.8635 | 0.8543 | 0.8734 | 0.8705 | 0.8512 | 0.8626 | 0.0087 |
| MAE (testset) | 0.6584 | 0.6560 | 0.6705 | 0.6659 | 0.6541 | 0.6610 | 0.0062 |
| Fit time | 1391.76 | 1385.70 | 1397.81 | 1403.85 | 1396.46 | 1395.12 | 6.09 |
| Test time | 15.34 | 16.92 | 13.63 | 10.31 | 13.91 | 14.02 | 2.19 |

```
In [49]: 1 # Printing out the results for these algorithms
2 print('Evaluation Results:')
3 print('Algorithm\t RMSE\t\t MAE')
4 print()
5 print('KNN Basic', '\t', round(cv_knn_basic['test_rmse'].mean(), 4), '\t')
6 print('KNN Means', '\t', round(cv_knn_means['test_rmse'].mean(), 4), '\t')
7 print('KNN ZScore', '\t', round(cv_knn_z['test_rmse'].mean(), 4), '\t',
8 print()
9 print('SVD', '\t\t', round(cv_svd['test_rmse'].mean(), 4), '\t', round(c
10 print('SVDpp', '\t\t', round(cv_svd_pp['test_rmse'].mean(), 4), '\t', ro
```

Evaluation Results:

| Algorithm | RMSE | MAE |
|------------|--------|--------|
| KNN Basic | 0.947 | 0.7254 |
| KNN Means | 0.8971 | 0.6852 |
| KNN ZScore | 0.8963 | 0.6798 |
| SVD | 0.8746 | 0.6723 |
| SVDpp | 0.8626 | 0.661 |

I chose two standard errors as our evaluation metrics:

- **Mean Absolute Error(MAE)** computes the average of all the absolute value differences between the true and the predicted rating.
- **Root Mean Square Error(RMSE)** computes the mean value of all the differences squared between the true and the predicted ratings and then proceeds to calculate the square root out of the result.

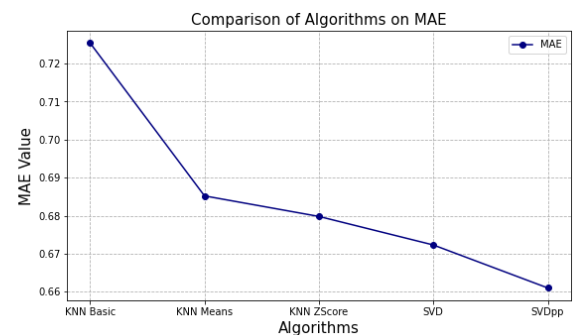
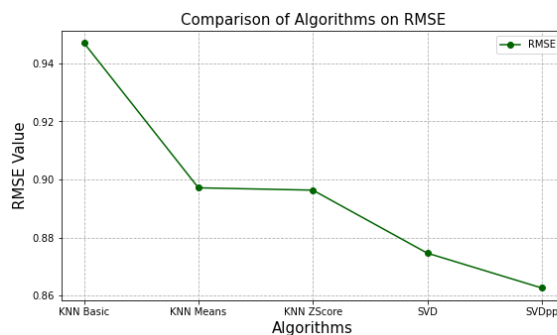
For both of these metrics, lower the error better the accuracy.

An average MAE of 0.6713 for SVD indicates an average absolute error of 0.6713 between the true and predicted ratings. I will try to reduce this error further by tuning hyperparameters.

```

In [50]: 1 # Plotting graphs for comparing accuracy of each algo
2 x_algo = ['KNN Basic', 'KNN Means', 'KNN ZScore', 'SVD', 'SVDpp',]
3 all_algos_cv = [cv_knn_basic, cv_knn_means, cv_knn_z, cv_svd, cv_svd_pp]
4
5 rmse_cv = [round(res['test_rmse'].mean(), 4) for res in all_algos_cv]
6 mae_cv = [round(res['test_mae'].mean(), 4) for res in all_algos_cv]
7
8 plt.figure(figsize=(20,5))
9
10 # RMSE graph
11 plt.subplot(1, 2, 1)
12 plt.title('Comparison of Algorithms on RMSE', loc='center', fontsize=15)
13 plt.plot(x_algo, rmse_cv, label='RMSE', color='darkgreen', marker='o')
14 plt.xlabel('Algorithms', fontsize=15)
15 plt.ylabel('RMSE Value', fontsize=15)
16 plt.legend()
17 plt.grid(ls='dashed')
18
19 # MAE graph
20 plt.subplot(1, 2, 2)
21 plt.title('Comparison of Algorithms on MAE', loc='center', fontsize=15)
22 plt.plot(x_algo, mae_cv, label='MAE', color='navy', marker='o')
23 plt.xlabel('Algorithms', fontsize=15)
24 plt.ylabel('MAE Value', fontsize=15)
25 plt.legend()
26 plt.grid(ls='dashed')
27 #plt.savefig('Images/RMSE_MAE')
28 plt.show()

```



Both the Matrix Factorization algorithms seem to do much better for both the metrics. KNN Means and KNN ZScore are also okay as compared to KNN Basic.

Selecting top models from each algorithm type: **KNN Means** and **SVDpp**

KNN Based Algorithms

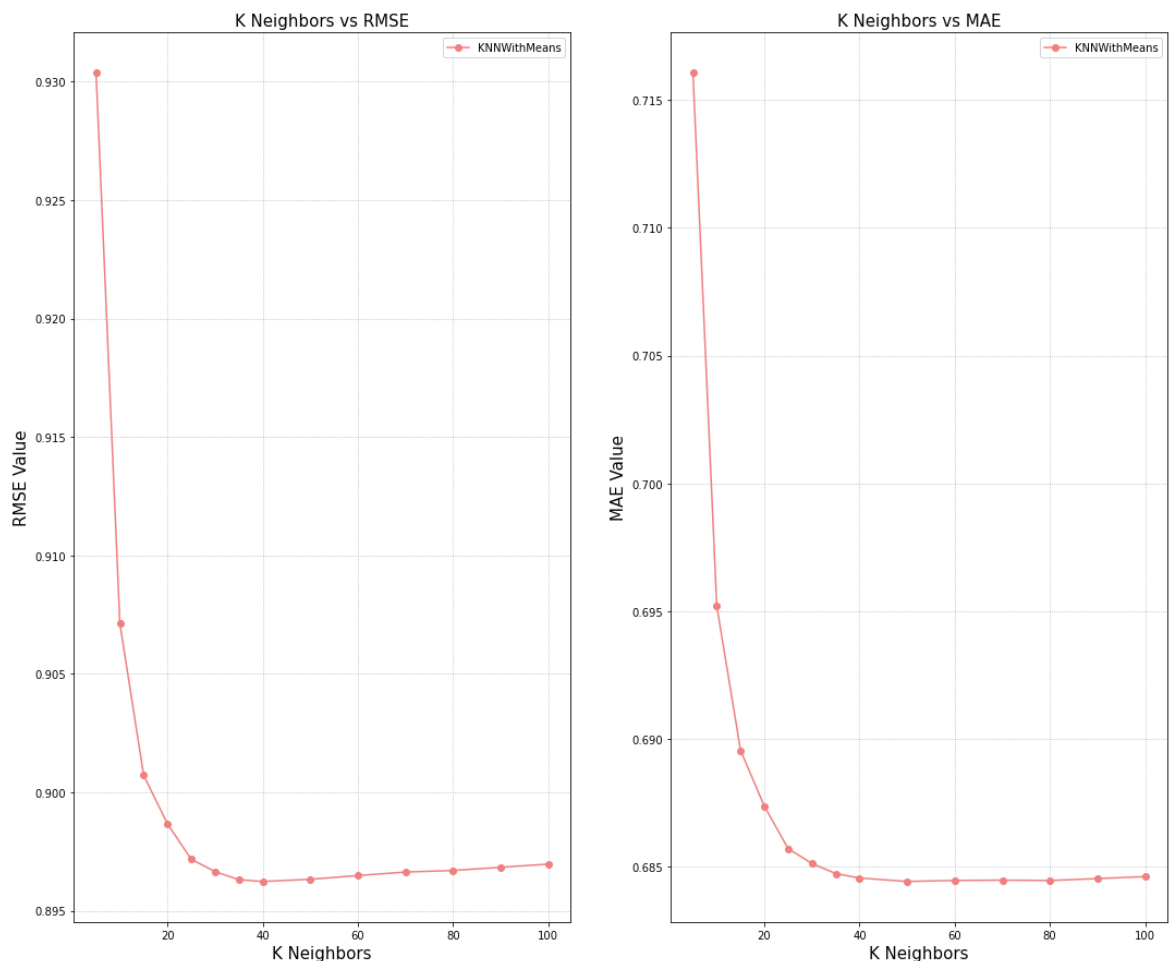
I will now optimize on these two models. Let's start with KNN Means. We will optimize two hyperparameters: k(number of neighbors and distance metric). First, I search for optimal k between 5 and 100.


```
In [51]: ▶ 1 param_grid = {'k': [5, 10, 15, 20, 25, 30, 35, 40, 50, 60, 70, 80, 90, 100]  
2  
3 gs_knn_means = GridSearchCV(KNNWithMeans, param_grid, measures=['rmse',  
4 gs_knn_means.fit(data)  
5  
6 y1 = gs_knn_means.cv_results['mean_test_rmse']  
7 y2 = gs_knn_means.cv_results['mean_test_mae']
```

```

In [52]: 1 # plotting accuracies for comparison
2 plt.figure(figsize = (18, 15))
3
4 x = [5, 10, 15, 20, 25, 30, 35, 40, 50, 60, 70, 80, 90, 100]
5
6 plt.subplot(1, 2, 1)
7 plt.title('K Neighbors vs RMSE', loc='center', fontsize=15)
8 plt.plot(x, y1, label='KNNWithMeans', color='lightcoral', marker='o')
9 plt.xlabel('K Neighbors', fontsize=15)
10 plt.ylabel('RMSE Value', fontsize=15)
11 plt.legend()
12 plt.grid(ls='dotted')
13
14 plt.subplot(1, 2, 2)
15 plt.title('K Neighbors vs MAE', loc='center', fontsize=15)
16 plt.plot(x, y2, label='KNNWithMeans', color='lightcoral', marker='o')
17 plt.xlabel('K Neighbors', fontsize=15)
18 plt.ylabel('MAE Value', fontsize=15)
19 plt.legend()
20 plt.grid(ls='dotted')
21 plt.savefig('Images/K-Neighbor vs RMSE')
22 plt.show()

```



From the plots above optimal k is found at 20. Now, we will look for the best distance metric out of these four: cosine, pearson, msd and pearson baseline.


```

In [53]: 1 # comparing distance matrix
2
3 knn_means_cosine = cross_validate(KNNWithMeans(k=20, sim_options={'name':
4 knn_means_pearson = cross_validate(KNNWithMeans(k=20, sim_options={'name'
5 knn_means_msd = cross_validate(KNNWithMeans(k=20, sim_options={'name': 'ms
6 knn_means_pearson_baseline = cross_validate(KNNWithMeans(k=20, sim_option
7
8
9 x_distance = ['cosine', 'pearson', 'msd', 'pearson_baseline',]
10 all_distances_cv = [knn_means_cosine, knn_means_pearson, knn_means_msd, k
11
12 rmse_cv = [round(res['test_rmse'].mean(), 4) for res in all_distances_cv]
13 mae_cv = [round(res['test_mae'].mean(), 4) for res in all_distances_cv]
14
15 plt.figure(figsize=(20,5))
16
17 plt.subplot(1, 2, 1)
18 plt.title('Comparison of Distance Metrics on RMSE', loc='center', fontsize
19 plt.plot(x_distance, rmse_cv, label='RMSE', color='darkgreen', marker='o'
20 plt.xlabel('Distance Metrics', fontsize=15)
21 plt.ylabel('RMSE Value', fontsize=15)
22 plt.legend()
23 plt.grid(ls='dashed')
24
25 plt.subplot(1, 2, 2)
26 plt.title('Comparison of Distance Metrics on MAE', loc='center', fontsize
27 plt.plot(x_distance, mae_cv, label='MAE', color='navy', marker='o')
28 plt.xlabel('Distance Metrics', fontsize=15)
29 plt.ylabel('MAE Value', fontsize=15)
30 plt.legend()
31 plt.grid(ls='dashed')
32 plt.savefig('Images/Comparison_of_Distance_metrics')
33 plt.show()

```

Evaluating RMSE, MAE of algorithm KNNWithMeans on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|-----|
| RMSE (testset) | 0.8928 | 0.9072 | 0.9052 | 0.9015 | 0.9120 | 0.9038 | 0.0 |
| 064 | | | | | | | |
| MAE (testset) | 0.6843 | 0.6960 | 0.6961 | 0.6877 | 0.6986 | 0.6925 | 0.0 |
| 055 | | | | | | | |
| Fit time | 1.04 | 1.10 | 1.24 | 1.17 | 1.24 | 1.16 | 0.0 |
| 8 | | | | | | | |
| Test time | 3.17 | 3.06 | 2.96 | 3.10 | 2.95 | 3.05 | 0.0 |
| 9 | | | | | | | |

Evaluating RMSE, MAE of algorithm KNNWithMeans on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|-----|
| RMSE (testset) | 0.9002 | 0.9050 | 0.8957 | 0.8942 | 0.9070 | 0.9004 | 0.0 |
| 050 | | | | | | | |
| MAE (testset) | 0.6864 | 0.6865 | 0.6817 | 0.6833 | 0.6917 | 0.6859 | 0.0 |
| 034 | | | | | | | |
| Fit time | 1.82 | 1.91 | 1.99 | 2.00 | 1.85 | 1.91 | 0.0 |
| 7 | | | | | | | |

```
Test time      4.08    3.99    3.86    3.85    3.84    3.93    0.0
9
```

Evaluating RMSE, MAE of algorithm KNNWithMeans on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|-----|
| RMSE (testset) | 0.8958 | 0.9097 | 0.8968 | 0.8881 | 0.9058 | 0.8992 | 0.0 |

| | | | | | | | |
|---------------|--------|--------|--------|--------|--------|--------|-----|
| MAE (testset) | 0.6827 | 0.6947 | 0.6871 | 0.6798 | 0.6939 | 0.6877 | 0.0 |
|---------------|--------|--------|--------|--------|--------|--------|-----|

| | | | | | | | |
|----------|------|------|------|------|------|------|-----|
| Fit time | 0.45 | 0.53 | 0.64 | 0.58 | 0.52 | 0.54 | 0.0 |
|----------|------|------|------|------|------|------|-----|

```
Test time      3.88    3.79    3.72    3.75    3.71    3.77    0.0
6
```

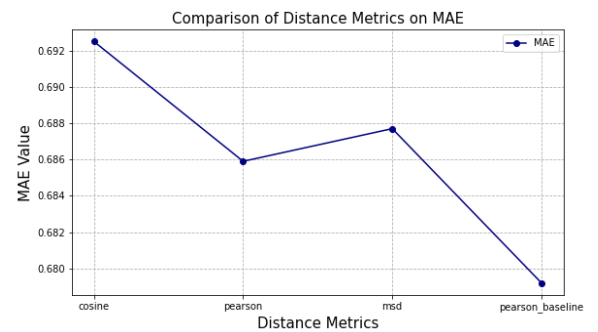
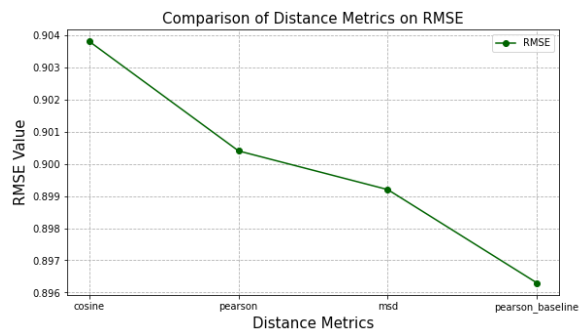
Evaluating RMSE, MAE of algorithm KNNWithMeans on 5 split(s).

| | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 | Mean | Std |
|----------------|--------|--------|--------|--------|--------|--------|-----|
| RMSE (testset) | 0.8917 | 0.8901 | 0.9006 | 0.8974 | 0.9015 | 0.8963 | 0.0 |

| | | | | | | | |
|---------------|--------|--------|--------|--------|--------|--------|-----|
| MAE (testset) | 0.6782 | 0.6770 | 0.6812 | 0.6775 | 0.6822 | 0.6792 | 0.0 |
|---------------|--------|--------|--------|--------|--------|--------|-----|

| | | | | | | | |
|----------|------|------|------|------|------|------|-----|
| Fit time | 1.78 | 1.89 | 1.79 | 1.89 | 1.68 | 1.80 | 0.0 |
|----------|------|------|------|------|------|------|-----|

```
Test time      3.67    3.71    3.69    3.61    3.57    3.65    0.0
5
```



Based on our hyperparameter tuning, the best KNN based model that we found out was:

KNN-Means with k=20 and pearson_baseline similarity

- **RMSE:** 0.8984
- **MAE :** 0.6807

Matrix Factorization Based Algorithms

While SVDpp had the better performance in terms of error rate, it is very time consuming to train. A grid search on SVDpp lasts many days. So, I chose to optimize the SVD model for number of epochs, learning rate and regularization using grid search.

```
In [54]: 1 #Parameter space
2 # run time approximately
3 svd_param_grid = {'n_epochs': [20, 25, 30, 40, 50],
4                     'lr_all': [0.007, 0.009, 0.01, 0.02],
5                     'reg_all': [0.02, 0.04, 0.1, 0.2]}
6
7 # This will take some time to execute.
8 gs_svd = GridSearchCV(SVD, svd_param_grid, measures=['rmse', 'mae'], cv=
9 gs_svd.fit(data)
```

```
In [55]: 1 print('Best value for SVD -RMSE:', round(gs_svd.best_score['rmse'], 4),
2 print('Optimal params RMSE =', gs_svd.best_params['rmse'])
3 print('optimal params MAE =', gs_svd.best_params['mae'])
```

Best value for SVD -RMSE: 0.8507 ; MAE: 0.6514

Optimal params RMSE = {'n_epochs': 50, 'lr_all': 0.01, 'reg_all': 0.1}

optimal params MAE = {'n_epochs': 50, 'lr_all': 0.01, 'reg_all': 0.1}

Based on our hyperparameter tuning, the best Matrix factorization based model that we found out was:

SVD with number of epochs = 50 , learning rate = 0.01 , and regularization = 0.1

- **RMSE:** 0.8507
- **MAE:** 0.6514

Predictions

Let's see the model in action to check if it's working as expected or not. But, first I need to fit both the final models on training set.

```
In [56]: 1 dataset = data.build_full_trainset()
2 final_knn_model = KNNWithMeans(k=20, sim_options={'name': 'pearson_base1
3 final_knn_model.fit(dataset)
```

Estimating biases using als...

Computing the pearson_baseline similarity matrix...

Done computing similarity matrix.

Out[56]: <surprise.prediction_algorithms.knns.KNNWithMeans at 0x233792abf40>

```
In [57]: 1 final_svd_model = SVD(n_epochs=50, lr_all=0.02, reg_all=0.1)
          2 final_svd_model.fit(dataset)
```

```
Out[57]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x23379486460>
```

Now that I have a trained model, I can use it to predict the rating a user would assign to a movie given an ID for the user (UID) and an ID for the item/movie(IID)

I am picking user with userID =610.This user really liked the movie Toy Story (1995) . We will now try to predict the rating that this user will give the movie Toy Story 3 (2010) .

```
In [58]: 1 # Showing rating given by this user for Toy Story (1995)
          2 # userId: 610
          3 # movieId: 1 for Toy Story (1995)
          4
          5 df[(df['userId'] == 610) & (df['movieId'] == 1)][['userId', 'movieId', 'rating', 'title']]
```

```
Out[58]:
```

| | userId | movieId | rating | title |
|-------|--------|---------|--------|------------------|
| 99534 | 610 | 1 | 5.0 | Toy Story (1995) |

Toy Story 3 (2010) was a well loved movie. We would assume that a user who really liked Toy Story 1 would really like this movie. So, we hope that the expected rating given by each of these models is quite high.

```
In [59]: 1 # KNN Model prediction on iid: 78499 - Toy Story 3 (2010)
          2 final_knn_model.predict(uid=610, iid=78499)
```

```
Out[59]: Prediction(uid=610, iid=78499, r_ui=None, est=4.646180057784888, details=
{'actual_k': 20, 'was_impossible': False})
```

```
In [60]: 1 # SVD Model prediction on iid: 78499 - Toy Story 3 (2010)
          2 final_svd_model.predict(uid=610, iid=78499)
```

```
Out[60]: Prediction(uid=610, iid=78499, r_ui=None, est=4.487028566763298, details=
{'was_impossible': False})
```

The field `est` indicates the estimated movie rating for this specific user.

We see that both the models give a strong positive rating for this specific movie and user. We also ran random experiments with a couple dozen user-movie pairs and received results that are consistent with our expectation.

We choose to move forward with the SVD model as it has a lower MAE and RMSE value.

We will now design a generic function which will take in a user id, then calculate expected ratings for all the movies and return the top 5 movies with the highest expected ratings.

We will also filter out movies already viewed by the user and provide functionality to mention preferred genre and minimum number of ratings.

```
In [61]: ▶ 1 def get_movie_recommendations(user_id, preferred_genre = 'all', minimum_n
2
3     new_df = df.copy()
4
5     # filtering out by genre
6     if preferred_genre != 'all':
7         new_df = new_df[new_df[preferred_genre]==1]
8
9     # filtering out by number of ratings
10    new_df = new_df[new_df['num rating']>=minimum_num_ratings]
11
12    # filtering out all movies already rated by user
13    movies_already_watched = set(new_df[new_df['userId']==user_id].movie
14    new_df= new_df[~new_df['movieId'].isin(movies_already_watched)]
15
16    # finding expected ratings for all remaining movies in the dataset
17    all_movie_ids = set(new_df['movieId'].values)
18    all_movie_ratings = []
19
20    for i in all_movie_ids:
21        expected_rating = final_svd_model.predict(uid=user_id, iid=i).es
22        all_movie_ratings.append((i,round(expected_rating,1)))
23
24    # extracting top five movies by expected rating
25    expected_df = pd.DataFrame(all_movie_ratings, columns=['movieId','Ex
26    result_df = pd.merge(expected_df, movies_df[['movieId','title','num
27    result_df = result_df.sort_values(['Expected Rating','num rating'],a
28
29    return result_df.head()
```

```
In [62]: ▶ 1 # receiving movie ratings for a given user id
2     get_movie_recommendations(1)
```

Out[62]:

| | movieId | Expected Rating | title | num rating |
|------------|---------|-----------------|---|------------|
| 73 | 318 | 5.0 | Shawshank Redemption, The (1994) | 317.0 |
| 174 | 858 | 5.0 | Godfather, The (1972) | 192.0 |
| 194 | 7153 | 5.0 | Lord of the Rings: The Return of the King, The... | 185.0 |
| 221 | 1221 | 5.0 | Godfather: Part II, The (1974) | 129.0 |
| 262 | 48516 | 5.0 | Departed, The (2006) | 107.0 |

Everything seems to be in place now. The above function utilizes both the model plus some filters

to give some truly amazing movie recommendations.

Minimum number of ratings is an interesting filter because if we set it too high, we only get classics and we won't find any new movies. Whereas if we set it too low, we can get virtually any movie. We like to think of it as an exploration risk parameter. Set value for it by asking yourself the following question: **How much risk are you willing to take to find new movies?**

Final Recommendation Engine

The final recommendation engine will be a hybrid between two models that we saw in this file: the final SVD model and the naive recommendation engine.

1. The naive recommendation engine is used to solve the **cold-start** problem for users who are new and have no ratings in the dataset.
2. If the userId is in the dataset, then we will use the final model with filters that we saw in the previous section.

```

In [63]: 1 def hybrid_recommendation_engine(user_id='new', preferred_genre='all', min
2
3     if user_id=='new':
4         if preferred_genre == 'all':
5             result = movie_ratings[(movies_df['num rating']>minimum_num_
6         else:
7             result = movie_ratings[(movies_df[preferred_genre] == 1) & (
8
9     else:
10        new_df = df.copy()
11
12        # filtering out by genre
13        if preferred_genre != 'all':
14            new_df = new_df[new_df[preferred_genre]==1]
15
16        # filtering out by number of ratings
17        new_df = new_df[new_df['num rating']>=minimum_num_ratings]
18
19        # filtering out all movies already rated by user
20        movies_already_watched = set(new_df[new_df['userId']==user_id].m
21        new_df = new_df[~new_df['movieId'].isin(movies_already_watched)]
22
23        # finding expected ratings for all remaining movies in the data
24        all_movie_ids = set(new_df['movieId'].values)
25        all_movie_ratings = []
26
27        for i in all_movie_ids:
28            expected_rating = final_svd_model.predict(uid=user_id, iid=i
29            all_movie_ratings.append((i, round(expected_rating, 1)))
30
31        # extracting top five movies by expected rating
32        expected_df = pd.DataFrame(all_movie_ratings, columns=['movieId'
33        result = pd.merge(expected_df, movies_df[['movieId', 'title', 'num
34        result = result.sort_values(['Expected Rating', 'num rating'], asc
35        result = result.head()
36
37
38        print('\n\nThese are the recommendations for the users with the follo
39        print('User id:', user_id)
40        print('Minimum number of ratings:', minimum_num_ratings)
41        print("User's choice of genre:", preferred_genre)
42        display(result)
43
44
45
46    genres = ['all',
47              'Animation',
48              'Children',
49              'Comedy',
50              'Fantasy',
51              'Romance',
52              'Drama',
53              'Action',
54              'Crime',
55              'Thriller',
56              'Horror',

```

```

57         'Mystery',
58         'Sci-Fi',
59         'War',
60         'Musical',
61         'Documentary',
62         'IMAX',
63         'Western',
64         'Film-Noir'
65     ]
66     all_userids = ['new'] + list(set(df.userId.values))
67     w = interactive(hybrid_recommendation_engine,
68                     user_id=widgets Dropdown(options=all_userids, description='User ID'),
69                     minimum_num_ratings=widgets.IntSlider(min=0, max=200, value=100, description='Minimum number of ratings'),
70                     preferred_genre=widgets Dropdown(options=genres, description='Preferred Genre'),
71                     )
72     display(w)

```

These are the recommendations for the users with the following filters
 User id: 2
 Minimum number of ratings: 100
 User's choice of genre: Animation

| | movieId | Expected Rating | title | num rating |
|---|---------|-----------------|-----------------------|------------|
| 9 | 68954 | 4.2 | Up (2009) | 105.0 |
| 3 | 6377 | 4.1 | Finding Nemo (2003) | 141.0 |
| 2 | 60069 | 4.1 | WALL·E (2008) | 104.0 |
| 4 | 364 | 4.0 | Lion King, The (1994) | 172.0 |
| 6 | 4306 | 4.0 | Shrek (2001) | 170.0 |

The hybrid recommendation engine allows us to effectively solve the cold-start problem and provide meaningful movie recommendations to all users.

Kindly try out the dashboard and let us know what you think. I have been using this for movie night recommendations.

Conclusion

I analyzed a variety of movie recommendation systems on the famous MovieLens database. I started with a naive recommendation engine which did not make any assumptions about the user

and provided general recommendations based upon movie popularity or the average ratings given by other users in the database.

I then progressed to some collaborative filtering based engines which try to find similar movies or users to make their predictions. After assessing models on two metrics, RMSE and MAE, we designed a SVD model and also tuned it for multiple hyperparameters.

Finally, I made a hybrid system of our naive recommendation engine and the SVD model to help resolve the cold-start problem. We added filtering options for genre and minimum number of ratings to give users some control over these recommendations.

Future Work

There is a lot of potential for future work in this project.

To begin with, I would like to add functionality in our final dashboard to allow new users to rate some movies and then to utilize that information to improve our recommendation system.

I also couldn't make use of `tag` information in this part of the analysis. We would like to make word embeddings from tags and other meta information about the movie and use it in our model.

I can also make use of the `links` dataset and scrape more information about each movie from the internet. This could involve significant features like cast, director, plot, etc.